

10-important Steps for Exploratory Data Analysis

1. Data Shape (columns or Rows ki taddad dekh len)
2. Check Data structure of each column or series
3. Missing values in each column and whole data set
4. Split variables or make new columns if needed
5. Type casting
6. Summary statistics
7. Value counts of a specific column
8. Deal with duplicates
9. Check the normal distribution of data (data Anomally)
10. Correlation between two variables (columns/series)

EDA karne se hamen kia kia maloom hta hy?

1. Data invetigation
2. Patterns inside the data
3. Anomalies (normal disribution hy ya skewed)
4. hyothsis konsa or kaisay design karna
5. Assumption konsi hni chahyeayn
6. Data visualization (Sirf pattern dekhney k liay)

```
In [ ]: #import Libararies
import numpy as np
import pandas as pd
import seaborn as sns
import scipy as sc
import matplotlib.pyplot as plt
```

```
In [ ]: #Load datasets
df= sns.load_dataset('tips')
df1= sns.load_dataset('titanic')
```

1.find shape

```
In [ ]: # shape

rows, cols = df.shape
print("Number of Rows      = ", rows) #instances
print("Number of Columns    = ", cols) #series
```

```
Number of Rows      = 244
Number of Columns    = 7
```

```
In [ ]: # shape

rows, cols = df1.shape
```

```
print("Number of Rows      = ", rows) #instances
print("Number of Columns   = ", cols) #series
```

```
Number of Rows      = 891
Number of Columns   = 15
```

2. Check Data structure of each column or series

In []:

```
# info of data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   total_bill  244 non-null    float64
1   tip         244 non-null    float64
2   sex         244 non-null    category
3   smoker      244 non-null    category
4   day         244 non-null    category
5   time        244 non-null    category
6   size        244 non-null    int64
dtypes: category(4), float64(2), int64(1)
memory usage: 7.4 KB
```

In []:

```
# data info
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   survived    891 non-null    int64
1   pclass      891 non-null    int64
2   sex         891 non-null    object
3   age         714 non-null    float64
4   sibsp       891 non-null    int64
5   parch       891 non-null    int64
6   fare        891 non-null    float64
7   embarked    889 non-null    object
8   class       891 non-null    category
9   who         891 non-null    object
10  adult_male   891 non-null    bool
11  deck         203 non-null    category
12  embark_town  889 non-null    object
13  alive        891 non-null    object
14  alone        891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
```

3. Missing values in each column and whole data set

In []:

```
# how many missing values present
df.isnull().sum()
```

```
Out[ ]: total_bill    0
tip              0
sex              0
smoker           0
```

```

day            0
time           0
size           0
dtype: int64

```

```

In [ ]: # how many missing values present
        df1.isnull().sum()

```

```

Out[ ]: survived            0
        pclass              0
        sex                 0
        age                 177
        sibsp               0
        parch               0
        fare                0
        embarked            2
        class               0
        who                 0
        adult_male          0
        deck                688
        embark_town         2
        alive               0
        alone               0
        dtype: int64

```

We will calculate the percentage of missing values, and if the percentage of missing value is high then we will reduce the priority of that column. We can deal with the missing values if the percentage is low by replacing them with the means, median or other methods (removal).

```

In [ ]: # percentage of missing values
        df.isnull().sum() / df.shape[0] *100

```

```

Out[ ]: total_bill         0.0
        tip                0.0
        sex                0.0
        smoker             0.0
        day                0.0
        time               0.0
        size               0.0
        dtype: float64

```

```

In [ ]: # percentage of missing values
        df1.isnull().sum() / df.shape[0] *100

```

```

Out[ ]: survived           0.000000
        pclass             0.000000
        sex                0.000000
        age                72.540984
        sibsp              0.000000
        parch              0.000000
        fare               0.000000
        embarked           0.819672
        class              0.000000
        who                0.000000
        adult_male         0.000000
        deck               281.967213
        embark_town        0.819672
        alive              0.000000
        alone              0.000000
        dtype: float64

```

In this example we will not consider the column name deck as the percentage of missing value is

quite high (77.22%).

4. Split variables or make new columns if needed

```
In [ ]: # making a dataframe for example using pandas library

df2 = pd.DataFrame(np.array([["Lahore, Pakistan", 67, 100], ["Beijing, China", 5, 6],
                             ], columns=['address', 'males', 'females']))

df2.head()
```

```
Out [ ]:
```

	address	males	females
0	Lahore, Pakistan	67	100
1	Beijing, China	5	6
2	berlin, Germany	8	9

```
In [ ]: # if we want to separate address into city and country columns we will split like th
df2[['city', 'country']] = df2['address'].str.split(',', expand=True)
#to see the results
df2.head()
```

```
Out [ ]:
```

	address	males	females	city	country
0	Lahore, Pakistan	67	100	Lahore	Pakistan
1	Beijing, China	5	6	Beijing	China
2	berlin, Germany	8	9	berlin	Germany

5. Type casting

```
In [ ]: # how to see the types in first place
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   address     3 non-null      object
1   males       3 non-null      object
2   females     3 non-null      object
3   city        3 non-null      object
4   country     3 non-null      object
dtypes: object(5)
memory usage: 248.0+ bytes
```

```
In [ ]: # convert data type into integer
df2[['males', 'females']] = df2[['males', 'females']].astype('int')
#convert to string
df2[['city', 'country']] = df2[['city', 'country']].astype('str').astype("string")
```

```
In [ ]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 3 entries, 0 to 2
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   address     3 non-null      object
 1   males       3 non-null      int32
 2   females     3 non-null      int32
 3   city        3 non-null      string
 4   country     3 non-null      string
dtypes: int32(2), object(1), string(2)
memory usage: 224.0+ bytes

```

```

In [ ]: # #Replace Data Types to Boolean
# df["IsPurchased"] = df['IsPurchased'].astype('bool')
# #Replace Data Types to Float
# df["Total Spend"] = df['Total Spend'].astype('float')
# #Replace Data Types to Datetime with format= '%Y%m%d'
# df['Dates'] = pd.to_datetime(df['Dates'], format='%Y%m%d')

```

6. Summary statistics

```

In [ ]: df.describe()

```

```

Out[ ]:

```

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672
std	8.902412	1.383638	0.951100
min	3.070000	1.000000	1.000000
25%	13.347500	2.000000	2.000000
50%	17.795000	2.900000	2.000000
75%	24.127500	3.562500	3.000000
max	50.810000	10.000000	6.000000

```

In [ ]: df1.describe()

```

```

Out[ ]:

```

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```

In [ ]: df2.describe()

```

```
Out[ ]:
```

	males	females
count	3.000000	3.000000
mean	26.666667	38.333333
std	34.961884	53.425961
min	5.000000	6.000000
25%	6.500000	7.500000
50%	8.000000	9.000000
75%	37.500000	54.500000
max	67.000000	100.000000

7. Value counts of a specific column

```
In [ ]: #how much values in a specific column
df1['age'].value_counts()
```

```
Out[ ]:
```

24.00	30
22.00	27
18.00	26
19.00	25
28.00	25
	..
36.50	1
55.50	1
0.92	1
23.50	1
74.00	1

Name: age, Length: 88, dtype: int64

```
In [ ]: df2['females'].value_counts()
```

```
Out[ ]:
```

100	1
6	1
9	1

Name: females, dtype: int64

```
In [ ]: df['tip'].value_counts()
```

```
Out[ ]:
```

2.00	33
3.00	23
4.00	12
5.00	10
2.50	10
	..
4.34	1
1.56	1
5.20	1
2.60	1
1.75	1

Name: tip, Length: 123, dtype: int64

```
In [ ]: #finding unique values in a column or series
df1['age'].unique()
```

```
Out[ ]: array([22. , 38. , 26. , 35. , nan, 54. , 2. , 27. , 14. ,
        4. , 58. , 20. , 39. , 55. , 31. , 34. , 15. , 28. ,
        8. , 19. , 40. , 66. , 42. , 21. , 18. , 3. , 7. ,
        49. , 29. , 65. , 28.5 , 5. , 11. , 45. , 17. , 32. ,
        16. , 25. , 0.83, 30. , 33. , 23. , 24. , 46. , 59. ,
        71. , 37. , 47. , 14.5 , 70.5 , 32.5 , 12. , 9. , 36.5 ,
        51. , 55.5 , 40.5 , 44. , 1. , 61. , 56. , 50. , 36. ,
        45.5 , 20.5 , 62. , 41. , 52. , 63. , 23.5 , 0.92, 43. ,
        60. , 10. , 64. , 13. , 48. , 0.75, 53. , 57. , 80. ,
        70. , 24.5 , 6. , 0.67, 30.5 , 0.42, 34.5 , 74. ])
```

```
In [ ]: df2['females'].unique()
```

```
Out[ ]: array([100, 6, 9])
```

8. Deal with duplicates

```
In [ ]: # find duplicates (remove)/ null values ( mean median mode)
df1[df1.embark_town == 'Queenstown']

#this will show the people only embarked from Queenstown in Titanic.
```

```
Out[ ]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	d
5	0	3	male	NaN	0	0	8.4583	Q	Third	man	True	1
16	0	3	male	2.0	4	1	29.1250	Q	Third	child	False	1
22	1	3	female	15.0	0	0	8.0292	Q	Third	child	False	1
28	1	3	female	NaN	0	0	7.8792	Q	Third	woman	False	1
32	1	3	female	NaN	0	0	7.7500	Q	Third	woman	False	1
...
790	0	3	male	NaN	0	0	7.7500	Q	Third	man	True	1
825	0	3	male	NaN	0	0	6.9500	Q	Third	man	True	1
828	1	3	male	NaN	0	0	7.7500	Q	Third	man	True	1
885	0	3	female	39.0	0	5	29.1250	Q	Third	woman	False	1
890	0	3	male	32.0	0	0	7.7500	Q	Third	man	True	1


77 rows × 15 columns



9. Check the normal distribution of data (data Anomally)


```
In [ ]: # plot histogram
sns.histplot(df1['age'])
```

```
Out[ ]: <AxesSubplot:xlabel='age', ylabel='Count'>
```



```
In [ ]: sns.histplot(df1['fare'])
```

```
Out[ ]: <AxesSubplot:xlabel='fare', ylabel='Count'>
```




This one is right skewed and doesn't normal at all

```
In [ ]: # Plotting 2 variables on the same graph

# set a grey background
sns.set(style="darkgrid")

sns.histplot(data=df1, x="age", color="skyblue", label="age", kde=True)
sns.histplot(data=df1, x="fare", color="red", label="fare", kde=True)

plt.legend()
plt.show()
```

The graph is looking messy because fare is not normally distributed and have outliers That's why it is better to look at them individually first.

```
In [ ]: df.head(2)
```

```
Out[ ]:
```

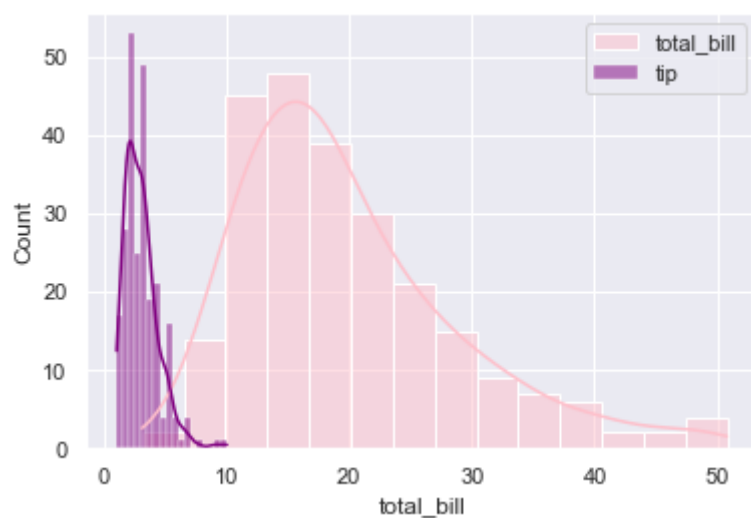
	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3

```
In [ ]: # Plotting 2 variables on the same graph

# set a grey background
sns.set(style="darkgrid")

sns.histplot(data=df, x="total_bill", color="pink", label="total_bill", kde=True)
sns.histplot(data=df, x="tip", color="purple", label="tip", kde=True)

plt.legend()
plt.show()
```




```
In [ ]: #measure its skewness and kurtosis
df1['age'].agg(['skew', 'kurtosis']).transpose()
```

```
skew    0.389108
```

```
Out[ ]: kurtosis    0.178274
        Name: age, dtype: float64
```

```
In [ ]: sns.histplot(df['tip'])
```

```
Out[ ]: <AxesSubplot:xlabel='tip', ylabel='Count'>
```



```
In [ ]: #measure its skewness and kurtosis
        df['tip'].agg(['skew', 'kurtosis']).transpose()
```


```
Out[ ]: skew        1.465451
        kurtosis     3.648376
        Name: tip, dtype: float64
```

```
In [ ]: sns.boxplot(df1['age'], color="red")
```

C:\Users\Azka\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[ ]: <AxesSubplot:xlabel='age'>
```




```
In [ ]: sns.boxplot(df['tip'], color="blue")
```

C:\Users\Azka\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning:

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
Out[ ]: <AxesSubplot:xlabel='tip'>
```



10. Correlation between two variables (columns/series)

Checking the correlation between variables is also necessary to see potential a feature that we can use for further analysis or building a model later. We can use a correlation matrix to get this.


```
In [ ]: # drawing correlation
corr = df.corr(method="pearson") # you can use spearman if you want
corr
# this will display a correlation matrix
```

```
Out[ ]:
```

	total_bill	tip	size
total_bill	1.000000	0.675734	0.598315
tip	0.675734	1.000000	0.489299
size	0.598315	0.489299	1.000000

```
In [ ]: sns.heatmap(corr, annot=True)
# this will show the numbers with colors
```

```
Out[ ]: <AxesSubplot:>
```



```
In [ ]: corr.style.background_gradient(cmap='gist_stern')
```

```
Out[ ]:
```

	total_bill	tip	size
total_bill	1.000000	0.675734	0.598315
tip	0.675734	1.000000	0.489299
size	0.598315	0.489299	1.000000


```
In [ ]: # drawing correlation
corr1 = df1.corr(method="pearson") # you can use spearman if you want
corr1
# this will display a correlation matrix
```

```
Out[ ]:
```

	survived	pclass	age	sibsp	parch	fare	adult_male	alone
survived	1.000000	-0.338481	-0.077221	-0.035322	0.081629	0.257307	-0.557080	-0.203367
pclass	-0.338481	1.000000	-0.369226	0.083081	0.018443	-0.549500	0.094035	0.135207
age	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	0.096067	0.280328	0.198270
sibsp	-0.035322	0.083081	-0.308247	1.000000	0.414838	0.159651	-0.253586	-0.584471
parch	0.081629	0.018443	-0.189119	0.414838	1.000000	0.216225	-0.349943	-0.583398
fare	0.257307	-0.549500	0.096067	0.159651	0.216225	1.000000	-0.182024	-0.271832
adult_male	-0.557080	0.094035	0.280328	-0.253586	-0.349943	-0.182024	1.000000	0.404744
alone	-0.203367	0.135207	0.198270	-0.584471	-0.583398	-0.271832	0.404744	1.000000

```
In [ ]: sns.heatmap(corr1, annot=True)
# this will show the numbers with colors
```

```
Out[ ]: <AxesSubplot:>
```



```
In [ ]: corr1.style.background_gradient(cmap='coolwarm')
```

```
Out[ ]:
```

	survived	pclass	age	sibsp	parch	fare	adult_male	alone
survived	1.000000	-0.338481	-0.077221	-0.035322	0.081629	0.257307	-0.557080	-0.203367
pclass	-0.338481	1.000000	-0.369226	0.083081	0.018443	-0.549500	0.094035	0.135207
age	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	0.096067	0.280328	0.198270
sibsp	-0.035322	0.083081	-0.308247	1.000000	0.414838	0.159651	-0.253586	-0.584471
parch	0.081629	0.018443	-0.189119	0.414838	1.000000	0.216225	-0.349943	-0.583398
fare	0.257307	-0.549500	0.096067	0.159651	0.216225	1.000000	-0.182024	-0.271832
adult_male	-0.557080	0.094035	0.280328	-0.253586	-0.349943	-0.182024	1.000000	0.404744
alone	-0.203367	0.135207	0.198270	-0.584471	-0.583398	-0.271832	0.404744	1.000000

```
In [ ]: #correlation between male vs female
df1male =df1[df1['sex']=='male']
df1female =df1[df1['sex']=='female']
```

```
In [ ]: # drawing correlation
corr1ma = df1male.corr(method="pearson")
corr1ma # this will display a correlation matrix
corr1ma.style.background_gradient(cmap='coolwarm')
```

```
Out[ ]:
```

	survived	pclass	age	sibsp	parch	fare	adult_male	alone
survived	1.000000	-0.220618	-0.119618	-0.020238	0.096318	0.171288	-0.234337	-0.133419
pclass	-0.220618	1.000000	-0.392754	0.076957	-0.031481	-0.472452	-0.078919	0.135319
age	-0.119618	-0.392754	1.000000	-0.334982	-0.232419	0.077331	0.536159	0.211858
sibsp	-0.020238	0.076957	-0.334982	1.000000	0.524849	0.181804	-0.461831	-0.637488
parch	0.096318	-0.031481	-0.232419	0.524849	1.000000	0.312197	-0.497120	-0.606243
fare	0.171288	-0.472452	0.077331	0.181804	0.312197	1.000000	-0.056082	-0.321650
adult_male	-0.234337	-0.078919	0.536159	-0.461831	-0.497120	-0.056082	1.000000	0.414375

	survived	pclass	age	sibsp	parch	fare	adult_male	alone
alone	-0.133419	0.135319	0.211858	-0.637488	-0.606243	-0.321650	0.414375	1.000000

```
In [ ]: # drawing correlation
corr1fe = df1male.corr(method="pearson")
corr1fe # this will display a correlation matrix
corr1fe.style.background_gradient(cmap='coolwarm')
```

	survived	pclass	age	sibsp	parch	fare	adult_male	alone
survived	1.000000	-0.220618	-0.119618	-0.020238	0.096318	0.171288	-0.234337	-0.133419
pclass	-0.220618	1.000000	-0.392754	0.076957	-0.031481	-0.472452	-0.078919	0.135319
age	-0.119618	-0.392754	1.000000	-0.334982	-0.232419	0.077331	0.536159	0.211858
sibsp	-0.020238	0.076957	-0.334982	1.000000	0.524849	0.181804	-0.461831	-0.637488
parch	0.096318	-0.031481	-0.232419	0.524849	1.000000	0.312197	-0.497120	-0.606243
fare	0.171288	-0.472452	0.077331	0.181804	0.312197	1.000000	-0.056082	-0.321650
adult_male	-0.234337	-0.078919	0.536159	-0.461831	-0.497120	-0.056082	1.000000	0.414375
alone	-0.133419	0.135319	0.211858	-0.637488	-0.606243	-0.321650	0.414375	1.000000

```
In [ ]: #calculate the correlation between the two arrays male vs female
np.corrcoef(corr1ma, corr1fe)[0,1]
```


Out[]: -0.37296113527583546

```
In [ ]: # drawing correlation
corr2 = df2.corr(method="pearson") # you can use spearman if you want
corr2
# this will display a correlation matrix
```

	males	females
males	1.00000	0.99989
females	0.99989	1.00000

```
In [ ]: sns.heatmap(corr2, annot=True)
# this will show the numbers with colors
```

Out[]: <AxesSubplot:>




```
In [ ]: corr2.style.background_gradient(cmap='YlOrBr_r')
```

```
Out[ ]:
```

	males	females
males	1.000000	0.999890
females	0.999890	1.000000

```
In [ ]: # we can also draw a pairplot to see the correlation
sns.pairplot(corr1)
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x1edc6ab5f40>
```



```
In [ ]: # we can also draw a pairplot to see the correlation
sns.pairplot(corr)
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x1edc9f473d0>
```

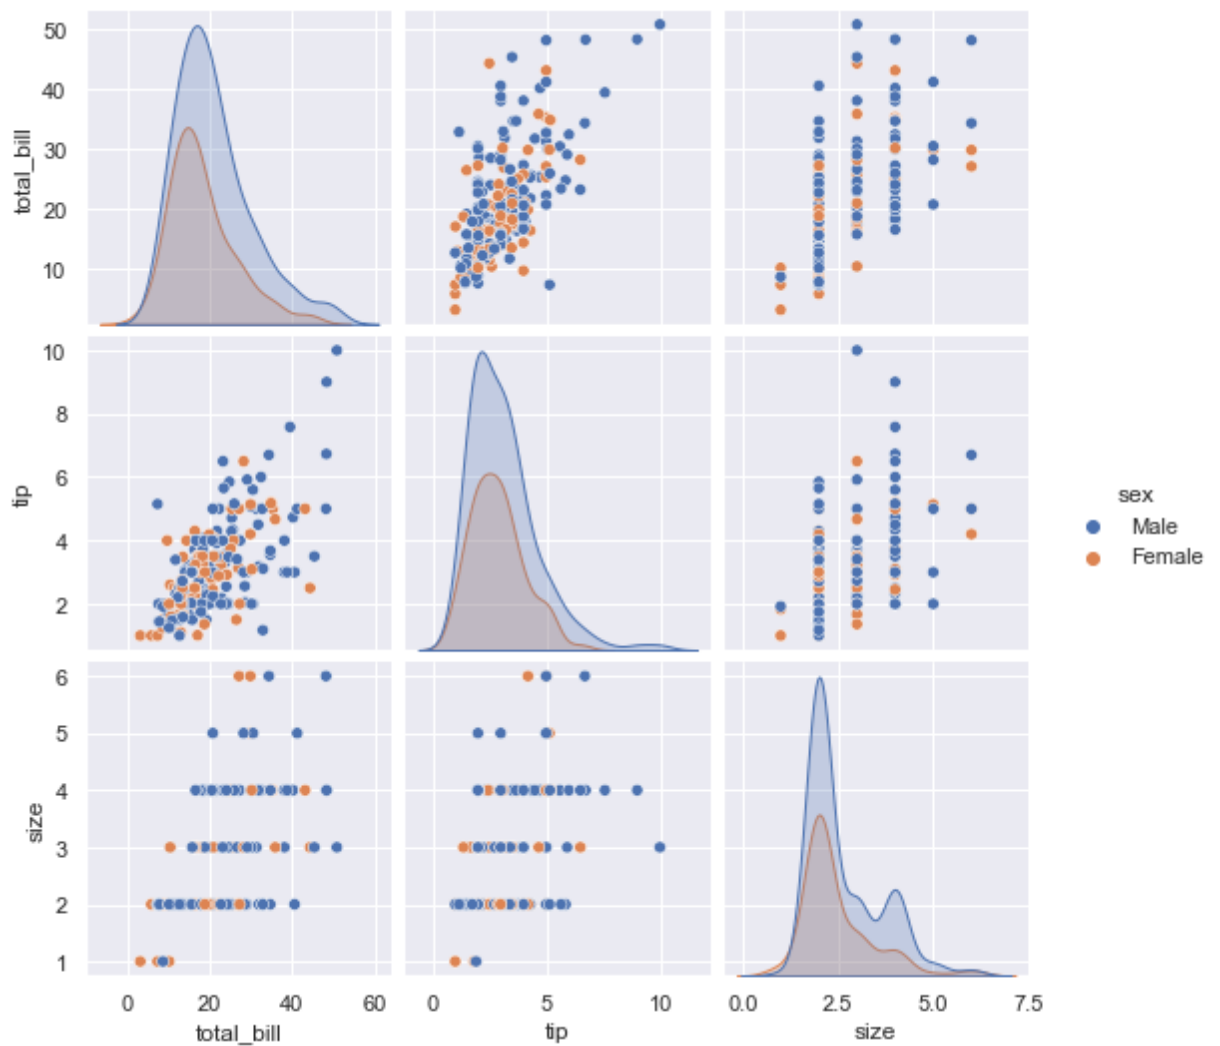



Out[]:

In []:


Out[]:

17/20



```
In [ ]: # we can convert this into histograms
sns.pairplot(df, hue="sex", diag_kind="hist")
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x1edc9f4e8e0>
```



```
In [ ]: # to make one sided
sns.pairplot(df, hue="sex", diag_kind="hist", corner=True)
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x1edcc9311c0>
```

